

Core DevOps Concepts

1. What is DevOps?

DevOps is a set of practices that automates the processes between software development and IT teams, in order that they can build, test, and release software faster and more reliably. It emphasizes collaboration and communication between development and operations teams.

2. What are the benefits of DevOps?

- Faster time to market
- Increased efficiency and productivity
- Improved collaboration and communication
- Higher quality software releases
- Increased customer satisfaction
- Reduced operational costs
- More stable operating environments

3. What are the key DevOps principles?

- Collaboration and communication
- Automation
- Continuous integration and continuous delivery (CI/CD)
- Infrastructure as code (IaC)
- Continuous monitoring and feedback
- Culture of shared responsibility

4. What is the difference between DevOps and Agile?

Agile is a software development methodology that focuses on iterative development and collaboration within development teams. DevOps extends Agile principles to the entire software lifecycle, including operations, by emphasizing automation, continuous delivery, and collaboration between development and operations.

5. What is Continuous Integration?

Continuous Integration (CI) is a development practice where developers regularly merge their code changes into a central repository, after which automated builds and tests are run. This helps detect integration errors early and ensures that the codebase remains stable.

6. What is Continuous Delivery?

Continuous Delivery (CD) is a software development practice where code changes are automatically built, tested, and prepared for release to production. It ensures that software can be released to production at any time with minimal effort.

7. What is Continuous Deployment?

Continuous Deployment is a software development practice where code changes that pass automated testing are automatically deployed to production. It takes Continuous Delivery a step further by automating the final deployment stage.

8. What is the difference between Continuous Delivery and Continuous Deployment?

Continuous Delivery means that code changes are automatically prepared for release, but a human must manually approve the release. Continuous Deployment automates the entire process, automatically deploying code changes to production after passing automated tests.

9. What are the biggest challenges in implementing DevOps?

- Cultural resistance to change
- Lack of collaboration and communication
- Legacy infrastructure and processes
- Security concerns
- Lack of skilled personnel
- Difficulty in automating everything
- Tool sprawl and integration issues

10. What are DevOps anti-patterns?

- Siloed teams (dev vs. ops)
- Manual deployments
- Ignoring feedback loops
- Over-reliance on tools without cultural change
- Treating DevOps as a role or team, not a culture
- Lack of automated testing
- Neglecting security (DevSecOps)
- Ignoring monitoring and observability
- "Throwing it over the wall" mentality
- Micromanagement of teams

Docker & Containers

11. What is Docker?

Docker is a platform for developing, shipping, and running applications in lightweight, portable containers. Containers allow developers to package an application with all of its dependencies into a standardized unit for software development.

12. What are the advantages of using Docker?

- Portability: Containers run consistently across different environments.
- Efficiency: Containers share the host OS kernel, reducing resource usage.
- Isolation: Containers isolate applications, preventing conflicts.
- Speed: Containers start quickly, improving deployment times.
- Scalability: Easy to scale applications by running multiple containers.
- Simplified Configuration: Dockerfiles provide a standardized way to configure environments.

13. What is the difference between Docker Image and Docker Container?

- Docker Image: A read-only template with instructions for creating a Docker container. It's like a blueprint.
- Docker Container: A runnable instance of a Docker image. It's the actual running application.

14. What is Dockerfile?

A Dockerfile is a text document that contains all the commands a user could call on the command line to assemble an image. It specifies the base image, adds application code, installs dependencies, and configures the environment.

15. How does Docker Compose work?

Docker Compose is a tool for defining and running multi-container Docker applications. It uses a YAML file (docker-compose.yml) to configure application services, networks, and volumes, allowing you to start and stop all services with a single command.

16. What is Docker Swarm?

Docker Swarm is Docker's native clustering and orchestration solution. It allows you to create and manage a cluster of Docker nodes (Swarm) and deploy services across the cluster.

17. What is the difference between Docker Swarm and Kubernetes?

- Kubernetes: More complex, feature-rich, and suitable for large, complex applications. Provides advanced orchestration features like auto-scaling, self-healing, and service discovery.
- Docker Swarm: Simpler, easier to set up, and suitable for smaller to medium-sized applications. Lacks some of the advanced features of Kubernetes.

18. Explain Docker Architecture.

Docker architecture consists of:

- Docker Client: The user interface to interact with Docker.
- Docker Daemon (dockerd): The background service that manages Docker images,

containers, networks, and volumes.

- Docker Registry: Stores Docker images (e.g., Docker Hub).
- Docker Images: Read-only templates for creating containers.
- Docker Containers: Runnable instances of Docker images.

19. What are Docker volumes?

Docker volumes are the preferred mechanism for persisting data generated by and used by Docker containers. They are managed by Docker and are stored outside the container's filesystem, allowing data to persist even if the container is deleted.

20. How do you optimize a Docker Image?

- Use a small base image.
- Minimize the number of layers (combine RUN commands).
- Remove unnecessary files and dependencies.
- Use multi-stage builds.
- Leverage Docker's build cache effectively.
- Use `.dockerignore` to exclude unnecessary files.
- Optimize application code and dependencies.

Kubernetes

21. What is Kubernetes?

Kubernetes (K8s) is an open-source container orchestration system for automating application deployment, scaling, and management. It provides a platform for automating the deployment, scaling, and operations of application containers across clusters of hosts.

22. What are the main components of Kubernetes architecture?

- Control Plane:
 - kube-apiserver: Exposes the Kubernetes API.
 - etcd: Consistent and highly-available key-value store used as Kubernetes' backing store for all cluster data.
 - kube-scheduler: Watches for newly created Pods with no assigned node and selects a node for them to run on.
 - kube-controller-manager: Runs controller processes.
 - cloud-controller-manager: Integrates with cloud provider APIs.
- Worker Nodes:
 - kubelet: An agent that runs on each node in the cluster. It makes sure that containers are running in a Pod.
 - kube-proxy: Maintains network rules on nodes.
 - Container Runtime (Docker, containerd, CRI-O): Runs containers.

23. What is a Pod in Kubernetes?

A Pod is the smallest deployable unit in Kubernetes. It represents a single instance of a running process in a cluster. A Pod can contain one or more containers that are tightly coupled and share the same network namespace and storage volumes.

24. What is a Service in Kubernetes?

A Service is an abstraction that defines a logical set of Pods and a policy by which to access them. It provides a stable network endpoint for accessing Pods, abstracting away the dynamic nature of Pod IPs.

25. What is a Namespace in Kubernetes?

A Namespace provides a mechanism for isolating groups of resources within a single cluster. It allows you to partition cluster resources among multiple users or teams.

26. What is a StatefulSet in Kubernetes?

A StatefulSet manages the deployment and scaling of a set of Pods, and provides guarantees about the ordering and uniqueness of these Pods. It is suitable for

applications that require stable network identifiers, persistent storage, and ordered deployment and scaling.

27. What is a DaemonSet in Kubernetes?

A DaemonSet ensures that all (or some) Nodes run a copy of a Pod. It is used to deploy system-level agents and services, such as log collectors and monitoring agents, that need to run on every node.

28. What is Helm?

Helm is a package manager for Kubernetes. It allows you to package, configure, and deploy applications and services onto Kubernetes clusters using Helm charts.

29. What is Kubernetes Ingress?

Ingress exposes HTTP and HTTPS routes from outside the cluster to Services within the cluster. It acts as a reverse proxy and load balancer, routing traffic to the appropriate Services based on hostnames and paths.

30. What are Kubernetes Operators?

Kubernetes Operators are applications that extend Kubernetes to automate the deployment, configuration, and management of complex applications. They encapsulate domain-specific knowledge and automate operational tasks, such as backups, upgrades, and scaling.

CI/CD (Continuous Integration & Continuous Deployment)

31. What is a CI/CD Pipeline?

A CI/CD pipeline is a series of automated steps that software undergoes from code commit to production deployment. It automates the build, test, and deployment processes, enabling faster and more reliable software releases.

32. What is Jenkins?

Jenkins is an open-source automation server that enables developers to automate the software build, test, and deployment processes. It is widely used for implementing CI/CD pipelines.

33. What are Jenkins Pipelines?

Jenkins Pipelines are a suite of plugins that support implementing and integrating continuous delivery pipelines into Jenkins. They allow you to define complex workflows as code, enabling automation and version control of your CI/CD processes.

34. What is GitLab CI?

GitLab CI (Continuous Integration) is a tool built into GitLab for software development using the continuous methodology. It allows you to automate the building, testing, and deployment of your applications directly from your GitLab repository.

35. What is CircleCI?

CircleCI is a cloud-based CI/CD platform that automates the build, test, and deployment of software. It integrates with GitHub, Bitbucket, and GitLab, and offers a user-friendly interface for configuring pipelines.

36. What is ArgoCD?

ArgoCD is a declarative, GitOps continuous delivery tool for Kubernetes. It automates the deployment of applications to Kubernetes based on configurations stored in Git repositories.

37. What is Tekton?

Tekton is a Kubernetes-native CI/CD framework that allows you to create and run CI/CD pipelines as Kubernetes resources. It provides a flexible and scalable way to automate software delivery.

38. What are Canary Deployments?

Canary deployments involve releasing a new version of an application to a small subset of users before rolling it out to the entire user base. This allows you to test the new version in a production environment and identify any issues before a full rollout.

39. What are Blue-Green Deployments?

Blue-green deployments involve running two identical production environments, one with the old version (blue) and one with the new version (green). Traffic is then switched from the blue environment to the green environment, minimizing downtime.

40. What are Rolling Deployments?

Rolling deployments involve gradually replacing old instances of an application with new instances. This is done by incrementally updating the application, ensuring minimal downtime and allowing for easy rollback if needed.

Cloud Platforms

41. What is Cloud Computing?

Cloud computing is the on-demand availability of computer system resources, especially data storage and computing power, without direct active management by the user. Large clouds often have functions distributed over multiple locations, each location being a data center.

42. What are the different types of cloud services?

- Infrastructure as a Service (IaaS)
- Platform as a Service (PaaS)
- Software as a Service (SaaS)
- Serverless Computing

43. What is AWS (Amazon Web Services)?

AWS is a comprehensive, evolving cloud computing platform provided by Amazon that includes a mixture of infrastructure as a service (IaaS), platform as a service (PaaS) and packaged software as a service (SaaS) offerings.

44. What is Azure?

Microsoft Azure is a cloud computing service created by Microsoft for building, testing, deploying, and managing applications and services through Microsoft-managed data centers.

45. What is Google Cloud Platform (GCP)?

Google Cloud Platform (GCP) is a suite of cloud computing services that runs on the same infrastructure that Google uses internally for its end-user products, such as Google Search, Gmail, and YouTube.

46. What is a Virtual Private Cloud (VPC)?

A Virtual Private Cloud (VPC) is a secure, isolated private cloud hosted within a public cloud. It allows you to launch AWS resources into a virtual network that you've defined.

47. What is Infrastructure as a Service (IaaS)?

IaaS provides you with access to computing resources—servers, storage, and networking—over the internet. You manage the operating systems, applications, and data, while the provider manages the underlying infrastructure.

48. What is Platform as a Service (PaaS)?

PaaS provides a platform allowing customers to develop, run, and manage applications without the complexity of building and maintaining the infrastructure typically associated with developing and launching an app. The provider manages the underlying infrastructure and operating systems.

49. What is Software as a Service (SaaS)?

SaaS is a software distribution model in which a third-party provider hosts applications and makes them available to customers over the internet. You access the software through a web browser or mobile app, without managing the underlying infrastructure or software.

50. What is Serverless Computing?

Serverless computing is a cloud computing execution model in which the cloud provider dynamically manages the allocation of machine resources. You build and run applications without managing servers. The provider handles scaling, patching, and infrastructure management.

Infrastructure as Code (IaC)

51. What is Infrastructure as Code?

Infrastructure as Code (IaC) is the process of managing and provisioning computer data centers through machine-readable definition files, rather than physical hardware configuration or interactive configuration tools.

52. What is Terraform?

Terraform is an open-source infrastructure as code tool that allows you to define and provision infrastructure using a declarative configuration language called HashiCorp Configuration Language (HCL). It supports multiple cloud providers and on-premises environments.

53. What is Ansible?

Ansible is an open-source automation tool that simplifies configuration management, application deployment, and task automation. It uses a simple, agentless architecture and YAML-based playbooks to define automation tasks.

54. What is the difference between Ansible and Terraform?

- Terraform: Focuses on provisioning and managing infrastructure resources (IaaS). It defines the desired state of infrastructure.
- Ansible: Focuses on configuring and managing software on existing infrastructure. It automates tasks and ensures systems are in a desired state.

55. What are Terraform providers?

Terraform providers are plugins that allow Terraform to interact with various cloud providers, SaaS providers, and other APIs. They define the resources and data sources that Terraform can manage.

56. What are Terraform modules?

Terraform modules are reusable configurations that encapsulate infrastructure components. They allow you to organize and reuse Terraform code, making it easier to manage complex infrastructure.

57. What is CloudFormation?

AWS CloudFormation is an Infrastructure as Code service that allows you to automate the provisioning and management of AWS infrastructure resources using JSON or YAML templates.

58. What is Pulumi?

Pulumi is an Infrastructure as Code tool that allows you to define and manage infrastructure using familiar programming languages like Python, JavaScript, TypeScript, Go, and .NET. It supports multiple cloud providers and on-premises environments.

59. What is Chef?

Chef is a configuration management tool that automates the configuration and management of infrastructure. It uses recipes and cookbooks to define the desired state of systems.

60. What is Puppet?

Puppet is a configuration management tool that allows you to automate the configuration and management of infrastructure. It uses a declarative language to define the desired state of systems.

Monitoring & Logging

61. What is monitoring in DevOps?

Monitoring in DevOps involves continuously collecting and analyzing data from infrastructure and applications to identify performance issues, detect anomalies, and ensure system health. It provides real-time visibility into the system's state, enabling proactive problem resolution.

62. What is ELK Stack?

ELK Stack is a collection of three open-source tools: Elasticsearch (search and analytics engine), Logstash (data processing pipeline), and Kibana (visualization and dashboard tool). It is used for log management and analysis.

63. What is Prometheus?

Prometheus is an open-source systems monitoring and alerting toolkit. It collects and stores time-series data, allowing you to query and visualize metrics.

64. What is Grafana?

Grafana is an open-source data visualization and monitoring tool. It allows you to create dashboards and visualizations from various data sources, including Prometheus, Elasticsearch, and others.

65. Explain the difference between monitoring and logging.

- Monitoring: Focuses on collecting and analyzing metrics (quantitative data) to track system performance and identify trends. It provides insights into the overall health and performance of the system.
- Logging: Focuses on collecting and analyzing logs (textual data) to understand application behavior, troubleshoot issues, and identify errors. It provides detailed information about events and activities within the system.

66. What is Fluentd?

Fluentd is an open-source data collector that allows you to unify data collection and consumption for better use and understanding of data. It is often used for log aggregation and forwarding.

67. What is OpenTelemetry?

OpenTelemetry is an open-source observability framework that provides APIs, SDKs, and tools for collecting telemetry data (metrics, logs, and traces). It aims to standardize the way telemetry data is generated and collected.

68. What is APM (Application Performance Monitoring)?

APM is the practice of monitoring and managing the performance and availability of software applications. It involves collecting and analyzing performance data to identify bottlenecks, troubleshoot issues, and optimize application performance.

69. What is log aggregation?

Log aggregation is the process of collecting logs from multiple sources and centralizing them in a single location for analysis. It simplifies log management and troubleshooting by providing a unified view of log data.

70. How do you set up centralized logging?

- Install and configure log collectors (e.g., Fluentd, Logstash) on servers and applications.
- Configure log collectors to forward logs to a centralized logging system (e.g., ELK Stack, Splunk).

- Set up a log storage and indexing system (e.g., Elasticsearch).
- Create dashboards and visualizations (e.g., Kibana, Grafana) to analyze log data.
- Implement alerting based on log patterns and errors.

Security & Compliance

71. What is DevSecOps?

DevSecOps is the practice of integrating security into every phase of the DevOps lifecycle, from planning to deployment and operations. It emphasizes shared responsibility for security among development, security, and operations teams.

72. What is Infrastructure Security?

Infrastructure security involves protecting the underlying hardware and software that support applications and services. This includes securing servers, networks, storage, and other infrastructure components from unauthorized access and attacks.

73. What is Container Security?

Container security involves protecting containerized applications and their underlying infrastructure from vulnerabilities and attacks. This includes securing container images, runtime environments, and orchestration platforms like Kubernetes.

74. What is Compliance as Code?

Compliance as Code is the practice of automating compliance checks and audits using code. It involves defining compliance requirements as code and integrating them into the CI/CD pipeline.

75. What are Security Best Practices in DevOps?

- Automate security testing in the CI/CD pipeline.
- Implement Infrastructure as Code (IaC) with security checks.
- Use container security scanning and vulnerability management.
- Implement access controls and least privilege principles.
- Use secrets management tools.
- Monitor and log security events.
- Perform regular security audits and penetration testing.
- Foster a security-aware culture.

76. What is Zero Trust Security?

Zero Trust Security is a security model that assumes no implicit trust within or outside an organization's perimeter. It requires strict identity verification for every user and device trying to access resources, regardless of their location.

77. What is SSL/TLS?

SSL (Secure Sockets Layer) and TLS (Transport Layer Security) are cryptographic protocols that provide secure communication over a computer network. They are used to encrypt data transmitted between a client and a server, ensuring confidentiality and integrity.

78. What is a Web Application Firewall (WAF)?

A WAF is a security device that monitors and filters HTTP traffic between a web application and the internet. It protects web applications from common web exploits, such as SQL injection and cross-site scripting (XSS).

79. What is Network Segmentation?

Network segmentation is the practice of dividing a network into smaller, isolated subnets. This limits the impact of security breaches and prevents attackers from moving laterally within the network.

80. What is Secrets Management?

Secrets management is the practice of securely storing and managing sensitive

information, such as passwords, API keys, and certificates. It involves using tools and techniques to prevent unauthorized access to secrets.

Linux Administration

81. What are the basic Linux commands every DevOps engineer should know?

- `ls`: List directory contents
- `cd`: Change directory
- `pwd`: Print working directory
- `mkdir`: Create directory
- `rm`: Remove files or directories
- `cp`: Copy files or directories
- `mv`: Move files or directories
- `cat`: Concatenate files and print on the standard output
- `grep`: Search for patterns in files
- `find`: Search for files in a directory hierarchy
- `chmod`: Change file permissions
- `chown`: Change file owner and group
- `ssh`: Secure shell login
- `scp`: Secure copy
- `ps`: Report a snapshot of current processes
- `top`/`htop`: Display Linux processes
- `netstat`/`ss`: Network statistics
- `ifconfig`/`ip`: Network interface configuration
- `systemctl`: Control the systemd system and service manager
- `journalctl`: Query the systemd journal
- `df`: Report file system disk space usage
- `du`: Estimate file space usage
- `tar`: Archive utility

82. What is Shell Scripting?

Shell scripting is writing scripts using a shell (like Bash) to automate tasks on a Unix-like operating system. These scripts can execute commands, manipulate files, and perform other operations.

83. What is systemd?

systemd is a system and service manager for Linux operating systems. It provides process management, service management, and system initialization capabilities. It replaces the older SysVinit system.

84. How do you manage services in Linux?

Using `systemctl`:

- `systemctl start <service>`: Start a service
- `systemctl stop <service>`: Stop a service
- `systemctl restart <service>`: Restart a service
- `systemctl status <service>`: Check service status
- `systemctl enable <service>`: Enable service to start on boot
- `systemctl disable <service>`: Disable service from starting on boot

85. What is Linux File System Hierarchy?

The Linux file system hierarchy defines the structure of directories and files in a Linux system. Key directories include:

- `/`: Root directory
- `/bin`: Essential user command binaries
- `/boot`: Static files of the boot loader
- `/dev`: Device files
- `/etc`: Host-specific system-wide configuration files
- `/home`: User home directories
- `/lib`: Essential shared libraries and kernel modules
- `/media`: Mount points for removable media
- `/mnt`: Mount points for temporarily mounted file systems
- `/opt`: Add-on application software packages

- `/proc`: Process information pseudo-filesystem
- `/root`: Root user's home directory
- `/run`: Runtime variable data
- `/sbin`: System administration binaries
- `/srv`: Site-specific data served by this system
- `/tmp`: Temporary files
- `/usr`: Secondary hierarchy for user programs
- `/var`: Variable files

86. How do you troubleshoot high CPU/memory usage in Linux?

- Use `top` or `htop` to identify processes consuming high CPU/memory.
- Use `ps` to get more detailed information about processes.
- Use `vmstat` and `free` to monitor memory usage.
- Use `strace` to trace system calls made by a process.
- Use `perf` to profile CPU usage.
- Check system logs (`/var/log/syslog`, `/var/log/messages`).
- Use `lsof` to identify open files and network connections.

87. What is process management in Linux?

Process management involves controlling and monitoring processes running on a Linux system. This includes starting, stopping, and monitoring processes, as well as managing process priorities and resource usage.

88. What is a cron job?

A cron job is a scheduled task in Linux that runs at specified intervals. It is managed by the cron daemon, which reads the crontab file to determine when to run tasks.

89. How do you manage permissions in Linux?

- `chmod`: Change file permissions (read, write, execute) for owner, group, and others.
- `chown`: Change file owner and group.
- `chgrp`: Change file group.
- Using numeric or symbolic notation (e.g., `chmod 755 <file>`, `chmod u+x <file>`).

90. What is SELinux?

SELinux (Security-Enhanced Linux) is a security architecture for Linux systems that provides mandatory access control (MAC). It enhances security by defining policies that control how processes and users interact with files and other system resources.

Version Control

91. What is Git?

Git is a distributed version control system that tracks changes in computer files and coordinates work among multiple people on those files. It allows for branching, merging, and reverting changes, enabling a collaborative and organized development process.

92. What is Git Branching Strategy?

A Git branching strategy is a set of rules and guidelines for how branches are used in a Git repository. It defines how features, releases, and hotfixes are managed using branches.

93. What is Git Flow?

Git Flow is a branching model designed around releases. It uses specific branches for different purposes, such as `master` for production releases, `develop` for integration, `feature` for new features, `release` for preparing releases, and `hotfix` for urgent fixes.

94. What is Trunk Based Development?

Trunk Based Development is a branching strategy where developers commit directly

to a single branch (trunk/main) and use short-lived feature branches for short-term work. It emphasizes continuous integration and frequent releases.

95. How to handle merge conflicts in Git?

- Identify the conflicting files.
- Open the files in a text editor and manually resolve the conflicts.
- Remove the conflict markers (<<<<<, =====, >>>>>).
- Stage the resolved files (`git add <file>`).
- Commit the merge (`git commit`).

96. What are Git hooks?

Git hooks are scripts that Git executes automatically before or after events such as commits, pushes, and receives. They can be used to automate tasks like code linting, testing, and enforcing commit message standards.

97. What is a Git submodule?

A Git submodule allows you to keep a Git repository as a subdirectory of another Git repository. It is used to include external dependencies or shared code within a project.

98. What is Git rebase vs merge?

- Git rebase: Integrates changes by moving commits from one branch onto another. It creates a linear commit history.
- Git merge: Integrates changes by creating a new merge commit that combines the changes from two branches. It preserves the commit history.

99. How do you squash commits in Git?

- Use `git rebase -i <commit>` to interactively rebase commits.
- Mark commits to be squashed with "s" or "squash".
- Edit the commit message to combine the squashed commits.
- Force push the rebased branch (`git push --force-with-lease`).

100. What is GitOps?

GitOps is a set of practices that uses Git as the single source of truth for declarative infrastructure and application configurations. It automates the deployment of changes to infrastructure and applications based on changes in Git repositories.

Scalability & High Availability

101. What is Scalability in DevOps?

Scalability in DevOps refers to the ability of a system to handle increased workload or demand by adding resources or capacity. It ensures that the system can maintain performance and availability as it grows.

102. What is High Availability?

High Availability (HA) refers to the ability of a system to remain operational and accessible even when components fail. It aims to minimize downtime and ensure continuous service availability.

103. What is Load Balancing?

Load balancing is the process of distributing network traffic or workload across multiple servers or resources. It improves performance, availability, and fault tolerance by preventing any single server from becoming overloaded.

104. What is Auto Scaling?

Auto Scaling is the ability of a system to automatically adjust its resources (e.g., compute instances) based on demand. It allows the system to scale up or down dynamically, ensuring optimal performance and cost efficiency.

105. What is Disaster Recovery?

Disaster Recovery (DR) is the process of recovering from a catastrophic event or

disaster that disrupts normal operations. It involves creating and implementing plans to restore data, systems, and services quickly and effectively.

106. What is multi-region deployment?

Multi-region deployment involves deploying applications and services across multiple geographical regions. It improves availability, fault tolerance, and performance by distributing workloads and data across different locations.

107. What is horizontal scaling vs vertical scaling?

- Horizontal scaling: Involves adding more machines or resources to a system to handle increased workload. It distributes the load across multiple instances.
- Vertical scaling: Involves increasing the resources (e.g., CPU, RAM) of a single machine to handle increased workload. It upgrades the existing hardware.

108. What is an active-active vs active-passive setup?

- Active-active setup: All instances of a system are actively processing traffic and handling requests simultaneously. It provides high availability and load balancing.
- Active-passive setup: One instance of a system is active and handling requests, while another instance is passive and standing by. If the active instance fails, the passive instance takes over.

109. What is database sharding?

Database sharding is the process of partitioning a large database into smaller, more manageable pieces called shards. Each shard contains a subset of the data and can be hosted on a separate server. It improves performance and scalability.

110. What is replication in databases?

Database replication is the process of copying data from one database server (master) to one or more other database servers (slaves). It provides data redundancy, improves read performance, and enables disaster recovery.

Backup & Disaster Recovery

111. What is Backup and Disaster Recovery?

Backup and Disaster Recovery (BDR) is a set of processes and tools used to protect data and systems from loss or damage and to restore them to a working state after a disruptive event. It ensures business continuity by minimizing downtime and data loss.

112. What are different types of backups?

- Full Backup: Copies all data to a backup medium.
- Incremental Backup: Copies only the data that has changed since the last backup (full or incremental).
- Differential Backup: Copies only the data that has changed since the last full backup.
- Snapshot Backup: Creates a point-in-time copy of data.
- Cloud Backup: Stores backup data in a cloud-based storage service.

113. What is RPO and RTO?

- RPO (Recovery Point Objective): The maximum acceptable amount of data loss measured in time. It determines how far back in time you need to recover data.
- RTO (Recovery Time Objective): The maximum acceptable downtime for an application or system. It determines how quickly you need to restore services.

114. What is Business Continuity Planning?

Business Continuity Planning (BCP) is the process of creating a plan to ensure that essential business functions can continue during and after a disaster. It involves identifying potential risks, developing recovery strategies, and testing the plan.

115. What are backup best practices?

- Implement a regular backup schedule.
- Use a combination of backup types (full, incremental, differential).
- Store backups in a secure and offsite location.
- Encrypt backup data.
- Test backups regularly to ensure they can be restored.
- Automate the backup process.
- Monitor backup jobs and logs.
- Document the backup and recovery procedures.
- Use versioning for backups.
- Implement a 3-2-1 backup strategy (3 copies of data, 2 different media, 1 offsite).

Cloud Native Architecture

116. What is Cloud Native Architecture?

Cloud Native Architecture is an approach to building and running applications that fully exploit the advantages of the cloud computing delivery model. It uses technologies like containers, microservices, service meshes, immutable infrastructure, and declarative APIs to enable loosely coupled systems that are resilient, manageable, and observable.

117. What are Microservices?

Microservices are an architectural approach that structures an application as a collection of small, independent services, built around business capabilities. Each service is independently deployable, scalable, and maintainable, enabling faster development cycles and improved fault isolation.

118. What is Service Mesh?

A Service Mesh is a dedicated infrastructure layer that makes service-to-service communication secure, fast, and reliable. It provides features like traffic management, security, and observability without requiring changes to application code. Examples include Istio and Linkerd.

119. What is Event-Driven Architecture?

Event-Driven Architecture (EDA) is a software architecture pattern centered around the concept of events. An event is a significant change in state. Systems react to these events by processing them, leading to decoupled and scalable systems.

120. What are the 12-Factor App principles?

The 12-Factor App is a methodology for building software-as-a-service applications that are:

- Codebase: One codebase tracked in revision control, many deploys.
- Dependencies: Explicitly declare and isolate dependencies.
- Config: Store config in the environment.
- Backing services: Treat backing services as attached resources.
- Build, release, run: Strictly separate build and run stages.
- Processes: Execute the app as one or more stateless processes.
- Port binding: Export services via port binding.
- Concurrency: Scale out via the process model.
- Disposability: Maximize robustness with fast startup and graceful shutdown.
- Dev/prod parity: Keep development, staging, and production as similar as possible.
- Logs: Treat logs as event streams.
- Admin processes: Run admin/management tasks as one-off processes.

API Gateway & Service Mesh

121. What is an API Gateway?

An API Gateway is a management tool that sits in front of an application programming interface (API) and acts as a single entry point for defined back-

end services. It handles tasks like routing, security, rate limiting, and monitoring.

122. What are the benefits of using API Gateway?

- Centralized Management: Provides a single point of entry for all API requests.
- Security: Handles authentication, authorization, and protection against common attacks.
- Routing: Routes requests to the appropriate backend services.
- Rate Limiting: Prevents abuse and ensures fair usage.
- Monitoring: Provides insights into API usage and performance.
- Transformation: Allows for request and response transformation.
- Caching: Improves performance by caching responses.

123. What is API Security?

API Security refers to the practices and technologies used to protect APIs from unauthorized access, attacks, and data breaches. It includes authentication, authorization, encryption, and protection against common API vulnerabilities.

124. What is Rate Limiting?

Rate limiting is a technique used to control the number of requests that a user or client can make to an API within a given time frame. It prevents abuse, ensures fair usage, and protects against denial-of-service attacks.

125. What is API Documentation?

API Documentation is a technical document that provides information about how to use an API. It includes details about endpoints, request/response formats, authentication, and examples, making it easier for developers to integrate with the API.

Cloud Cost Optimization

126. What is Cloud Cost Optimization?

Cloud Cost Optimization is the process of reducing overall cloud spending while maintaining or improving performance. It involves identifying and eliminating unnecessary costs, right-sizing resources, and leveraging cost-effective cloud services.

127. What are Reserved Instances?

Reserved Instances (RIs) are a pricing model offered by cloud providers like AWS, Azure, and GCP that provide significant discounts compared to on-demand pricing. In exchange for a one-time or monthly upfront payment, you commit to using a specific instance type for a defined term (e.g., 1 or 3 years).

128. What is Spot Instance pricing?

Spot Instance pricing is a pricing model that allows you to bid on unused compute capacity in the cloud. Spot Instances are typically offered at a significant discount compared to on-demand pricing, but they can be interrupted with little to no notice if the cloud provider needs the capacity back.

129. How to implement cost tagging strategy?

- Define a consistent tagging schema (e.g., department, project, environment, owner).
- Implement automated tagging policies.
- Educate teams on tagging best practices.
- Regularly audit and enforce tagging compliance.
- Use tags for cost allocation and reporting.
- Consider using tools to automate tag management.

130. What are cost allocation reports?

Cost allocation reports are reports that break down cloud spending based on specific criteria, such as tags, services, or accounts. They provide insights into where cloud costs are coming from, allowing you to identify areas for

optimization and allocate costs to the appropriate teams or projects.

Site Reliability Engineering (SRE)

131. What is Site Reliability Engineering?

Site Reliability Engineering (SRE) is a discipline that applies software engineering principles to IT operations. It focuses on automating and improving the reliability, scalability, and performance of systems, using metrics and data-driven approaches.

132. What are Service Level Objectives (SLOs)?

Service Level Objectives (SLOs) are specific, measurable goals for the reliability and performance of a service. They define the desired level of service and are typically expressed as a percentage over a specific time period (e.g., 99.9% uptime per month).

133. What are Service Level Indicators (SLIs)?

Service Level Indicators (SLIs) are metrics that measure the performance of a service. They quantify aspects of service behavior, such as latency, error rate, and throughput, and are used to determine if SLOs are being met.

134. What is Error Budget?

Error Budget is the amount of allowable downtime or service degradation within a given time period, based on the SLO. It represents the difference between the target SLO and the actual performance of the service. Teams can use the error budget to balance feature development with reliability.

135. What is Toil in SRE?

Toil in SRE refers to manual, repetitive, automatable, tactical, devoid of enduring value, and that scales linearly with service growth. It's the work that keeps the lights on but doesn't contribute to long-term improvements or innovation. SRE aims to minimize toil through automation and process improvements.

DevOps Metrics & KPIs

136. What are DevOps Metrics?

DevOps Metrics are quantifiable measures used to assess the effectiveness and efficiency of DevOps practices. They provide insights into the performance of development and operations processes, helping teams identify areas for improvement.

137. What is Mean Time to Recovery (MTTR)?

Mean Time to Recovery (MTTR) is the average time it takes to restore a failed service or system to a working state. It measures the speed and efficiency of incident response and recovery.

138. What is Change Failure Rate?

Change Failure Rate is the percentage of changes to production that result in incidents or failures. It measures the stability and reliability of the deployment process.

139. What is Deployment Frequency?

Deployment Frequency is the number of times code is deployed to production within a given time period. It measures the speed and agility of the deployment process.

140. What is Lead Time for Changes?

Lead Time for Changes is the time it takes for a code change to go from commit to production. It measures the efficiency and speed of the entire software delivery pipeline.

Incident Management

141. What is Incident Management?

Incident Management is the process of restoring normal service operation as quickly as possible and minimizing the adverse impact on business operations, thus ensuring that the best possible levels of service quality and availability are maintained.

142. What is an Incident Response Plan?

An Incident Response Plan is a documented set of procedures that outlines how an organization will respond to and manage an incident, minimizing disruption and ensuring a swift recovery.

143. What is Post-Mortem Analysis?

Post-Mortem Analysis is a detailed review of an incident after it has been resolved, focusing on understanding the root cause, identifying areas for improvement, and preventing future occurrences.

144. What are Incident Severity Levels?

Incident Severity Levels are classifications of incidents based on their impact on business operations, typically ranging from critical (highest impact) to minor (lowest impact).

145. What is On-Call Management?

On-Call Management is the process of organizing and managing the availability of personnel to respond to incidents outside of normal working hours, ensuring timely resolution of issues.

146. What is Incident Escalation?

Incident Escalation is the process of transferring an incident to a higher level of support or management when it cannot be resolved within a specified timeframe or requires additional expertise.

147. What is Root Cause Analysis (RCA) in Incident Management?

Root Cause Analysis (RCA) in Incident Management is the process of identifying the underlying cause(s) of an incident to prevent recurrence and improve system reliability.

148. What is a War Room in Incident Response?

A War Room in Incident Response is a centralized location (physical or virtual) where incident response teams gather to coordinate and manage the resolution of critical incidents.

149. What tools are commonly used for Incident Management?

Common tools include Jira Service Management, PagerDuty, ServiceNow, Opsgenie, and Statuspage.

150. What is Mean Time to Detect (MTTD)?

Mean Time to Detect (MTTD) is the average time it takes to identify that an incident has occurred.

DevOps Culture & Best Practices

151. What is DevOps Culture?

DevOps Culture emphasizes collaboration, communication, automation, and continuous improvement between development and operations teams, fostering a shared responsibility for the entire software lifecycle.

152. How do you implement a blameless postmortem?

Implement a blameless postmortem by focusing on understanding what happened and

why, without assigning blame to individuals. Encourage open communication and learning from mistakes.

153. What is Knowledge Sharing in DevOps?

Knowledge Sharing in DevOps involves disseminating information, best practices, and lessons learned across teams to improve collaboration, efficiency, and continuous learning.

154. How does DevOps improve developer productivity?

DevOps improves developer productivity by automating repetitive tasks, providing faster feedback loops, enabling continuous delivery, and fostering a collaborative environment.

155. What are the key components of a high-performing DevOps team?

Key components include cross-functional collaboration, automation, continuous integration and delivery, feedback loops, and a culture of continuous improvement.

Infrastructure Monitoring & Observability

156. What is Observability in DevOps?

Observability in DevOps is the ability to understand the internal state of a system based on its external outputs, enabling proactive troubleshooting and performance optimization.

157. What is the difference between Monitoring and Observability?

Monitoring tells you *if* something is wrong, while Observability tells you *why* something is wrong. Monitoring is about known unknowns, Observability is about unknown unknowns.

158. What are the three pillars of Observability?

The three pillars of Observability are metrics, logs, and traces.

159. What is Distributed Tracing?

Distributed Tracing is a technique used to track requests as they propagate through a distributed system, providing insights into latency, errors, and dependencies.

160. How do you set up alerting in a monitoring system?

Set up alerting by defining thresholds for key metrics, configuring notification channels (email, SMS, etc.), and implementing escalation policies.

Cloud Migration & Modernization

161. What is Cloud Migration?

Cloud Migration is the process of moving applications, data, and IT resources from on-premises infrastructure to a cloud computing environment.

162. What are the key Cloud Migration Strategies?

Key strategies include rehosting (lift and shift), replatforming, refactoring, rearchitecting, and replacing.

163. What is the Lift and Shift approach?

The Lift and Shift approach involves moving applications and data to the cloud without making significant changes to their architecture or code.

164. What is Application Modernization?

Application Modernization is the process of updating legacy applications to take advantage of modern technologies and cloud-native architectures.

165. What is a Hybrid Cloud Strategy?

A Hybrid Cloud Strategy involves using a combination of on-premises infrastructure and cloud services to meet business needs.

Performance Testing & Optimization

166. What is Performance Testing?

Performance Testing is the process of evaluating the speed, responsiveness, and stability of an application or system under various workloads.

167. What are the different types of Performance Tests?

Types include load testing, stress testing, endurance testing, spike testing, and scalability testing.

168. What are the best tools for Performance Testing?

Tools include JMeter, Gatling, LoadRunner, and k6.

169. How do you analyze Performance Test results?

Analyze results by examining metrics such as response time, throughput, error rate, and resource utilization to identify bottlenecks and areas for optimization.

170. What is Load Testing vs. Stress Testing?

Load Testing evaluates performance under expected workloads, while Stress Testing evaluates performance under extreme workloads to identify breaking points.

API Gateway & Service Mesh

171. What is a Service Mesh?

A Service Mesh is a dedicated infrastructure layer that handles service-to-service communication, providing features like traffic management, security, and observability.

172. How does Istio work?

Istio works by deploying sidecar proxies (Envoy) alongside application containers to intercept and manage traffic, providing a transparent layer for service communication.

173. What is Linkerd?

Linkerd is a lightweight and ultra-fast service mesh that provides observability, security, and reliability for microservices.

174. How does an API Gateway improve security?

An API Gateway improves security by handling authentication, authorization, rate limiting, and protection against common API vulnerabilities.

175. What is Rate Limiting in APIs?

Rate Limiting in APIs controls the number of requests a client can make within a given time frame to prevent abuse and ensure fair usage.

Advanced Container Orchestration

176. What is a Kubernetes Custom Resource Definition (CRD)?

A Kubernetes Custom Resource Definition (CRD) allows you to extend the Kubernetes API by defining custom resources and controllers.

177. What is Kubernetes Horizontal Pod Autoscaler (HPA)?

Kubernetes Horizontal Pod Autoscaler (HPA) automatically scales the number of

Pods in a deployment based on CPU utilization or custom metrics.

178. What is Kubernetes Vertical Pod Autoscaler (VPA)?

Kubernetes Vertical Pod Autoscaler (VPA) automatically adjusts the CPU and memory requests of Pods to optimize resource utilization.

179. What is Kubernetes Cluster Autoscaler?

Kubernetes Cluster Autoscaler automatically adjusts the size of the Kubernetes cluster by adding or removing nodes based on Pod resource requests.

180. What is a Sidecar Pattern in Kubernetes?

A Sidecar Pattern in Kubernetes involves deploying a utility container alongside the main application container to provide supporting functionalities like logging, monitoring, or security.

Cloud Cost Optimization (Advanced)

181. What is FinOps?

FinOps is the practice of bringing financial accountability to the variable spend model of cloud, enabling teams to make informed decisions about cloud usage.

182. How does Spot Instance pricing work?

Spot Instance pricing allows you to bid on unused compute capacity, but instances can be interrupted with short notice when capacity is needed back.

183. What are Cloud Cost Allocation Reports?

Cloud Cost Allocation Reports break down cloud spending based on tags, services, or accounts, providing insights for cost optimization.

184. How do you optimize storage costs in the cloud?

Optimize storage costs by using appropriate storage classes, deleting unused data, and leveraging data lifecycle policies.

185. What are some real-world cost-saving techniques for Kubernetes clusters?

Techniques include right-sizing nodes, using spot instances, implementing resource quotas, and leveraging cluster autoscaling.

Advanced DevOps Metrics & KPIs

186. What is Mean Time Between Failures (MTBF)?

Mean Time Between Failures (MTBF) is the average time between system failures, indicating reliability.

187. What is Change Lead Time?

Change Lead Time is the time it takes for a code change to go from commit to production.

188. What is Deployment Success Rate?

Deployment Success Rate is the percentage of successful deployments out of total deployments.

189. How do you measure DevOps ROI?

Measure DevOps ROI by analyzing metrics like deployment frequency, lead time, MTTR, and change failure rate, and correlating them with business outcomes.

190. What are the four key DORA Metrics?

The four key DORA Metrics are Deployment Frequency, Lead Time for Changes, Mean Time to Recovery (MTTR), and Change Failure Rate.

Serverless Architecture (Advanced)

191. What is Function Cold Start in Serverless?

Function Cold Start in Serverless is the delay that occurs when a serverless function is invoked for the first time or after a period of inactivity. It's the time taken to provision and initialize the execution environment.

192. How do you reduce latency in a Serverless function?

- Keep function code and dependencies lightweight.
- Use provisioned concurrency or keep-alive mechanisms (where available).
- Minimize external dependencies and network calls.
- Optimize database queries and connections.
- Choose a region geographically close to users.
- Use caching mechanisms.

193. What is an Event-Driven Serverless Architecture?

An Event-Driven Serverless Architecture is a design pattern where serverless functions are triggered by events, such as file uploads, database changes, or messages from queues. It promotes loose coupling and scalability.

194. What is the difference between FaaS and PaaS?

- FaaS (Function as a Service): Focuses on running individual functions in response to events, with the provider managing the underlying infrastructure.
- PaaS (Platform as a Service): Provides a platform for developing, running, and managing applications, with the provider managing infrastructure and runtime environments.

195. How do you handle database connections in a Serverless environment?

- Use connection pooling to reuse database connections.
- Leverage serverless-optimized database services (e.g., AWS Aurora Serverless).
- Minimize the number of database connections per function invocation.
- Use environment variables or secrets management to store database credentials.
- Handle connection timeouts and retries gracefully.

Advanced DevSecOps & Security

196. What is Infrastructure Drift?

Infrastructure Drift is the divergence between the desired state of infrastructure (as defined in code or configuration) and the actual state of the deployed infrastructure. It occurs when manual changes or external factors alter the infrastructure without updating the configuration.

197. How do you implement secrets management in DevOps?

- Use dedicated secrets management tools (e.g., HashiCorp Vault, AWS Secrets Manager, Azure Key Vault).
- Store secrets encrypted and separate from application code.
- Implement least privilege access to secrets.
- Rotate secrets regularly.
- Avoid hardcoding secrets in configuration files or code.
- Integrate secrets management with CI/CD pipelines and runtime environments.

198. What is Role-Based Access Control (RBAC)?

Role-Based Access Control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within an organization. It assigns permissions to roles and then assigns roles to users.

199. How does DevSecOps handle compliance requirements?

- Automate compliance checks using Compliance as Code.
- Integrate security and compliance testing into the CI/CD pipeline.
- Use infrastructure as code (IaC) to ensure consistent and auditable

configurations.

- Implement continuous monitoring and logging for compliance reporting.
- Conduct regular security audits and vulnerability assessments.
- Foster a security-aware culture across all teams.

200. What are the key OWASP Top 10 security risks for DevOps?

- Injection (e.g., SQL injection, command injection)
- Broken Authentication
- Sensitive Data Exposure
- XML External Entities (XXE)
- Broken Access Control
- Security Misconfiguration
- Cross-Site Scripting (XSS)
- Insecure Deserialization
- Using Components with Known Vulnerabilities
- Insufficient Logging & Monitoring